CS4225/CS5425 Big Data Systems for Data Science
Course Project Guidelines

1. Overview

This is a team project. Each team comprises 2 or 3 members (CS4225 students can form a team with CS5425 students and vice versa). Each team should identify a topic on Big Data Systems for Data Science, and conduct a course project specified in one of the topics in the "Project Topics" section. You are encouraged to utilize other systems learnt from the lecture in data science such as those in data visualization and machine learning.

An important component of the course is the course project. Rather than having more homework, the course project allows students the flexibility to explore some aspect of reconfigurable computing on their own. The project involves the development of a proposal, experimental results, and a final report. In proposing a project please keep the following in mind: it is better to propose a reasonable project you can complete rather than a huge project that will be in an intermediate state at the end of the term. To receive a final grade for the project, there must be an experimental result that is shown (e.g. a graph, table, or chart) based on experiments that you perform.

You are encouraged to use a cloud platform for processing big data. Microsoft has kindly offered us the education grant on Windows Azure for this course. Therefore, you may be able to run your project and analytics on multiple virtual machines on Azure. If you are interested in using those resources, please contact our teaching assistant.

2. Project schedule

10 Feb 2019	Grouping on IVLE
9 March 2019	Submit a project proposal (See Section 6, upload to IVLE)
16 April 2019	Project presentation
20 Apr 2019	Project due – report must be submitted (upload to IVLE)

* all deadlines are 11:59pm of the date.

3. Project Topics

The topic of the course project can be ONE of the followings:

- **Data Analysis**

In big data processing, data analysis is a very critical step since it draws conclusions from new datasets that are important to a specific domain. Nowadays, trending application domain contains healthcare, insurance, transportation, social media, etc. As a part of this project, you need to identify: a) A trending topic that is interesting to you. b) Novel and influential analytic methods that are relevant to this topic. (You need to find some papers) c) New datasets that are not fully explored by data scientists. (*Better if the datasets are released within 1 year*).

Impact: you may create new data science insights that can impact the society and improve people's life. For example, by analyzing the taxi trajectory data carefully, you may help to cut down the waiting time for each user and make the trip more environment friendly.

Examples: a) A real-time sentiment analysis of twitter feeds with the NASDAQ index (i.e., analyze the correlation between tweeter feeds and hourly movements of the NASDAQ index), b) Use deep learning techniques to predict the stock price.

- **Scalable Data Science Tools and Platforms**

Scalable platforms and tools are very important for meeting the requirement of big data. Although there are many new algorithms that have been introduced, they may work on a single machine. In this project, we will develop scalable platforms and tools and make them open-sourced for public use. Here, scalability can be across multiple cores or multiple machines. Examples include the new data mining and machine learning algorithms that was proposed recently.

Impact: you may create new platforms and tools that can address the big data challenge that have not been ever attacked. For example, by developing an easy-to-use and efficient graph processing platform (say finding the shortest path), you may help the graph analysts to improve their productivity.

Examples: a) Visualized graph analytics: the study implements different ways of graph visualization and offers an intuitive way of graph analytics. b) K-means optimization: the study optimizes K-means on multi-core machines, and later on a distributed environment. It studies different implementation strategies (e.g., early reduction) and hardware capabilities (e.g., Ethernet vs. RDMA).

- **Special topics from NUS-Grab joint lab**

Supported by NUS-Grab joint lab (http://ids.nus.edu.sg/Grab-NUS-AI-Lab.html), we will create the opportunities for five teams to develop systems and applications on the data sets from Grab. The quota will be filled in the first-come-first-serve manner by sending the lecturer an email to indicate the interest. The lab has one-month data set for Grab car and taxi in Singapore. The data set includes: 1) the trip information that you can see from the Grab app, such as source and destination, discount and purpose of the trip. 2) the location information of the car/taxi at per second basis. You can choose one of the topics below. All the implementations are based on Apache Spark. Since the course project will be conducted on the real data set, team members have to be officially affiliated with the joint lab, and sign the necessary NDA document.

- Topic 1: Traffic flow analysis. Using the per-second location information of car/taxi, we want to analyse and predict the traffic flow of a city. Can we predict the speed of each road? It is very helpful to a lot of decision making operations if this prediction can be made timely and precisely.
  Ref.: Pablo Samuel Castro, Daqing Zhang, and Shijian Li. 2012. Urban traffic modelling and prediction using large scale taxi GPS traces. In Proceedings of the 10th international conference on Pervasive Computing (Pervasive'12). https://pdfs.semanticscholar.org/2990/2cfc29c3fd20e1a5c52d1e86058aa8908b7b.pdf.
- Topic 2: Point-of-interest analysis and prediction. Given a particular area (e.g., The Merlion) of the city, we want to know the top-k trajectories that will pass by the point-of-interest at certain time period (say 12pm to 1pm). This is helpful in managing traffic condition around that particular area.
  Ref.: Shuyao Qi, Panagiotis Bouros, Dimitris Sacharidis, Nikos Mamoulis. Efficient Point-Based Trajectory Search. SSTD 2015: 179-196. https://i2.cs.hku.hk/~nikos/SSTD15.pd
- Topic 3: Similar trajectories search. How to quickly exact similar trajectories given all driver spatial and temporal location information. This is a fundamental and challenging operation that is useful for further traffic and transportation optimisation analysis.
  Ref.: Dong Xie, Feifei Li, and Jeff M. Phillips. 2017. Distributed trajectory similarity search. Proc. VLDB Endow. 10, 11 (August 2017), 1478-1489. https://www.cs.utah.edu/~lifeifei/papers/trajvldb17.pdf.
- Topic 4: Trajectory compression. There are many existing algorithms (see Ref) for compressing trajectory data. A practical question is whether they can be efficiently applied to real dataset.
  Ref.: Dongxiang Zhang, Mengting Ding, Dingyu Yang, Yi Liu, Ju Fan, and Heng Tao Shen. 2018. Trajectory simplification: an experimental study and quality analysis. Proc. VLDB Endow. 11, 9 (May 2018), 934-946. http://www.vldb.org/pvldb/vol11/p934-zhang.pdf.

- **Self-proposed Projects**

If you have some ideas that may not fit exactly what is listed above (e.g., you want to design a totally new algorithm, or new statistical measures of interestingness, etc), talk to the lecturer.

Impact: The only limit to your impact is your imagination and commitment. (by Tony Robbins).

Examples: a new machine learning platform.

4. More Example Data Sets and Sample Projects

You can use any datasets on the internet. We prefer to use the **new and big** data sets (*Better if the datasets are released within 1 year*). There are several possible URL links:

https://www.kaggle.com/datasets

https://cloud.google.com/public-datasets/

https://github.com/caesar0301/awesome-public-datasets

https://github.com/openimages/dataset

http://www.cs.cmu.edu/~enron/

https://webobservatory.soton.ac.uk/

……

Regarding sample big data projects, you can refer to the following links. However, we encourage you to think beyond that and explore significantly new ideas (for example, whether you should identify a trending topic that is interesting to you, and then apply novel and influential analytic methods that are relevant to this topic).

https://www.coursera.org/specializations/big-data

http://hadoopproject.com/big-data-projects/

https://blog.kaspersky.com/cool-big-data-projects/8186/

……

5. Project Proposal

The purpose of the project proposal is to provide background material for the work that you are to complete and to describe the actual experiments and expected results. The proposal should be sufficiently detailed so that I can understand specifically what you are going to do. An important part of the proposal is my understanding of your proposed work. If I think the project is too large I will ask you to trim it down. The following is a general outline for the proposal. *All page counts are for 11pt. font, single spaced, single column.*

1) Topic introduction (0.5 - 1 page)
2) Discussion of previous work and how it relates to your topic. (DO NOT CUT AND PASTE PICTURES OR TEXT FROM OTHER DOCUMENTS IN COMPLETING THE PREVIOUS WORK SECTION) (1-2 pages)
3) Discussion of your experimental approach (including specific tools, methodologies, experiments, etc.). Try to be as specific as possible. (1 to 2 pages)
4) Data sets to be used (Are they large and new? Justify it, 0.5-1 pages)
5) Expected results. What exactly will your result be? What will you show in a table or chart? Be specific. (1 page)
6) Project Summary (0.5 to 1 page)
7) References (should include research articles and other sources, which is more than just web pages)

You need to submit your project proposal to IVLE, and name of the file should be include all student IDs of your team. For example, if the group has three members, the file name should be:

[GroupID]-[Student1ID]-[Student2ID]-[Student3ID]-proposal.pdf

6. Submission Requirements

a) Final Report

You need to submit a report which is at most 20 single-columned page paper on the problem and solution(s) and what will be demonstrated. *All page counts are for 11pt. font, single spaced, single column.*

You should compare the different solutions qualitatively as well as provide an experimental analysis. In general, the final report should contain the following contents:1. Project Introduction. This can be the same as for the proposal. 2. Methodology and experimentation. How did you perform the experimentation? This can be a revision of the proposal. 3. Discussion of results. 4. Problems encountered and lesson learnt. 5. Personal contribution (for each student written by the individual student). 6. Project summary. 7. References.

b) Code

Code can be in any programming languages (c, c++, java, matlab, R…).

c) File Name

You need to compress all the documents and code into a zip (or rar) file, and name of the file should be include all student IDs of your team. For example, if the group has three members, the file name should be:

[GroupID]-[Student1ID]- [Student2ID]-[Student3ID]-FinalReport.zip

Note: Please ensure that the code and report are complete before submission! You need to ensure that your work is recoverable by others using the code and report that you provided.

7. Submission Policies

For both project proposal and final report:

No multiple submissions allowed: Each team should make sure that the team only submit exactly once. If a team submits two versions, we will **retain the earliest version and discard all later versions**. The team will then be grade on the earliest version. For such a reason, if you want to update your submission, you should delete your old submission first and then submit a new one.

Policy on late submission: For fairness, reports submitted after the deadline but no more than 48 hours after the deadline will still be graded, with a penalty of 20%. Namely, I will first grade the report normally, and then multiple the mark by 80% to get the final mark for that report. **Reports submitted more than 48 hours after the deadline will not be accepted and will get 0 mark.**

8. Presentation

The presentations will take place during lecture. Your presentation should cover the same aspects of the project mentioned above. The final report will be due one week after oral presentation. Thus, you are encouraged to add more solid results and findings into the report, besides those presented in oral presentation.

9. Grading

Although students performing different works, to make the grading work fair and reasonable, we will evaluate your work from the following perspectives.

(a) Linguistic ability
(b) Complexity and novelty of the problem (you must carefully review the previous studies on the same problem)
(c) Tools and algorithms
(d) Comprehensiveness of the analysis and findings
(e) Impact of your project or findings
(f) Datasets used (Large? New?)
(g) Oral presentation and demo

10. Plagiarism

You are reminded that plagiarism is a very **SERIOUS** offense, and disciplinary action (including possibility of expulsion from the university) will be taken against any individual or team found plagiarizing. The individual or team that is being plagiarized will also be punished if it is found to have allowed the work to be plagiarized voluntarily.